

Introduction to Hardware Programming



NOTES:

QNX, Momentics, Neutrino and Photon microGUI are registered trademarks in certain jurisdictions, and PhAB, Phindows, and Qnet are trademarks of QNX Software Systems.

All other trademarks and trade names belong to their respective owners.

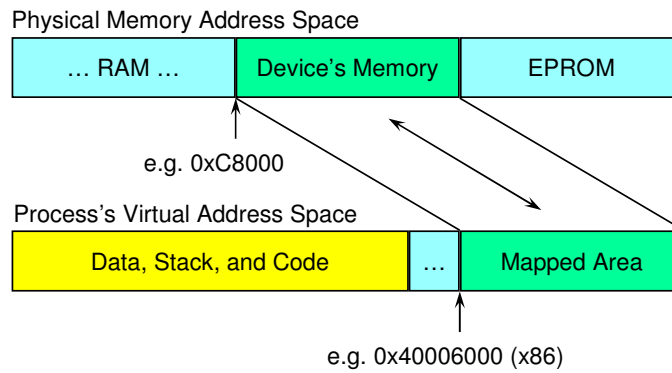
Topics:

- **Hardware I/O**
- Programming PCI bus devices**
- Handling Interrupts**
- Conclusion**

NOTES:

To access memory on a hardware device:

- physical addresses must be mapped into your process's virtual address space:



```
// get a pointer to memory physically located at 0xc8000
vaddr = mmap_device_memory (0, 0x4000,
                             PROT_READ | PROT_WRITE | PROT_NOCACHE,
                             0, 0xc8000);
```

NOTES:

DMA operations usually require physically contiguous RAM:

```
// for DMA:
// allocate len bytes of physically contiguous
// system memory

vaddr = mmap (0, len,
              PROT_READ | PROT_WRITE | PROT_NOCACHE,
              MAP_PHYS | MAP_ANON | MAP_PRIVATE, NOFD, 0);

// get the physical address for passing to the controller
mem_offset64 (vaddr, NOFD, len, &paddr, NULL);
```



NOTES:

For cards on an ISA bus, you may need the additional **MAP_BELOW16M** and **MAP_NOX64K** flags.

MAP_BELOW16M means do not go above 16M physical address.

MAP_NOX64K tells the process manager that the memory it allocates must not cross 64k boundary (physical address), useful on x86 only.

Accessing hardware registers:

```
// enable I/O privilege for this thread
ThreadCtl (_NTO_TCTL_IO, NULL);

// get access to a devices registers
iobase = mmap_device_io (len, base_port);

val8 = in8 (iobase+N);      // read an 8 bit value
val16 = in16 (iobase+N);    // read a 16 bit value
val32 = in32 (iobase+N);    // read a 32 bit value

out8 (iobase+N, val8);      // write an 8 bit value
out16 (iobase+N, val16);    // write a 16 bit value
out32 (iobase+N, val32);    // write a 32 bit value

- Include header: <hw/inout.h>
```



NOTES:

The *in*()* and *out*()* functions are done as inline functions, see `<$PROCESSOR/inout.h>` for details.

Topics:

Hardware I/O

→ **Programming PCI bus devices**

Handling Interrupts

Conclusion

NOTES:

To find and configure a PCI device:

- you must run one of the pci servers:

`pci-bios`, `pci-p5064`, ...

- you need to connect to the pci server with `pci_attach()` before making any other `pci_*`() calls

NOTES:

The PCI calls include:

<code>pci_attach()</code>	connect to PCI server, will fail if no PCI bus or server not running
<code>pci_detach()</code>	disconnect from PCI server
<code>pci_find_device()</code>	find hardware by Device ID and Vendor ID
<code>pci_find_class()</code>	find hardware by class
<code>pci_attach_device()</code>	find hardware and get basic configuration information
<code>pci_detach_device()</code>	release device configuration
<code>pci_read_config()</code>	read configuration information
<code>pci_read_config*()</code>	read blocks of 8/16/32-bit values
<code>pci_write_config()</code>	write configuration information
<code>pci_write_config*()</code>	write blocks of 8/16/32-bit values

NOTES:

The *pci_attach_device()* call:

- fills in a **pci_dev_info** structure
- the structure contains information about the PCI device, including:
 - **Irq** – interrupt number
 - an array of 6 memory areas:
 - **BaseAddressSize[i]** – size of area (0 for not used)
 - **CpuBaseAddress[i]** – address on the CPU side
 - **PciBaseAddress[i]** – address on the PCI side
 - **CpuBmstrTranslation** – address translations for bus master PCI devices

NOTES:

Topics:

Hardware I/O

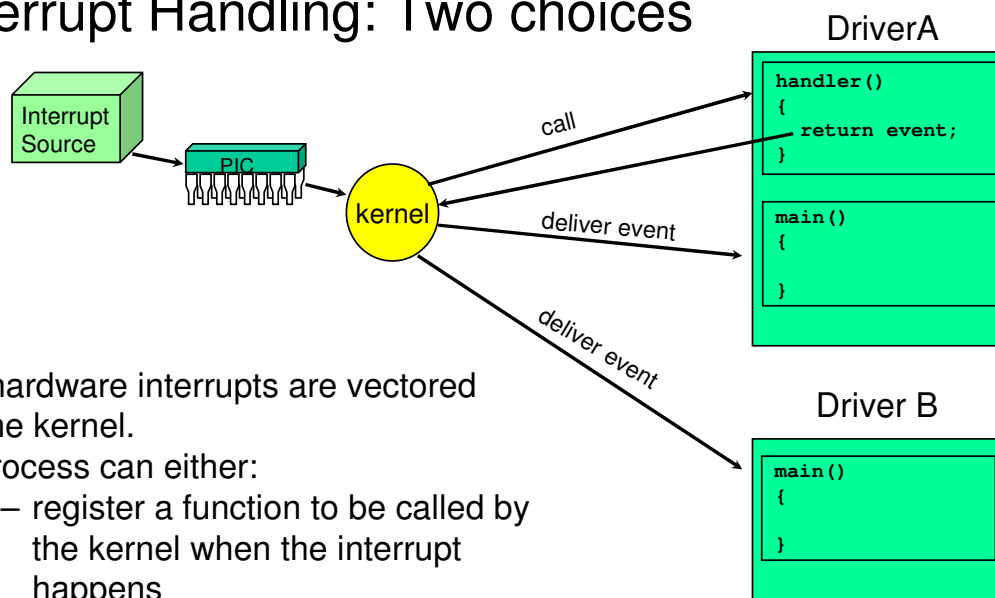
Programming PCI bus devices

→ Handling Interrupts

Conclusion

NOTES:

Interrupt Handling: Two choices



All hardware interrupts are vectored to the kernel.

A process can either:

- register a function to be called by the kernel when the interrupt happens
- request notification that the interrupt has happened

NOTES:

Interrupt calls:

```
id = InterruptAttach (int intr,
                    struct sigevent *(*handler)(void *, int),
                    void *area, int size, unsigned flags);
id = InterruptAttachEvent (int intr, struct sigevent *event,
                        unsigned flags);

InterruptDetach (int id);
InterruptWait (int flags, uint64_t *reserved);
InterruptMask (int intr, int id);
InterruptUnmask (int intr, int id);
InterruptLock (struct intrspin *spinlock);
InterruptUnlock (struct intrspin *spinlock);
```

☞ You must have I/O privilege for the above functions to work. To get I/O privilege, you call *ThreadCtl(_NTO_TCTL_IO, 0)*, and you must have root (userid 0) permissions.

NOTES:

There are also:

```
InterruptEnable (void);
InterruptDisable (void);
```

They are for non-SMP systems only and so are not good as a general solution. You should use *InterruptLock()* and *InterruptUnlock()* instead.

Driver A example: interrupt handler function

```
struct sigevent event;

const struct sigevent *
handler (void *not_used, int id)
{
    if (check_status_register())
        return (&event);
    else
        return (NULL);
}

main ()
{
    ThreadCtl (_NTO_TCTL_IO, 0);
    SIGEV_INTR_INIT (&event);
    id = InterruptAttach (intnum, handler, NULL, 0, ...);
    for (;;) {
        InterruptWait (0, NULL);
        // do some or all of the work here
    }
}
```

NOTES:

In this case, there is a user-supplied interrupt handler. This would be the case where there is work that must be done at interrupt priority for timing reasons. Notice that the handler and the thread can share the work. The handler can do the time-critical work (typically I/O) and the thread can be woken up every now and then for the non-time-critical work (number crunching, passing the data on to other threads.)

In the interrupt handler above, *check_status_register()* represents some code that checks some status register on the hardware to see if our hardware generated the interrupt and/or to clear the source of the interrupt. This is needed on level-sensitive architectures, since interrupts can be shared, and the kernel will issue an EOI at the end of the interrupt chain, so we must clear the interrupt before returning. This is also why, using the *InterruptAttachEvent()* method, the kernel must mask the interrupt before scheduling the thread.

Driver B example: interrupt event loop

```
struct sigevent event;

main ()
{
    ThreadCtl (_NTO_TCTL_IO, 0);
    SIGEV_INTR_INIT (&event);
    id = InterruptAttachEvent (intnum, &event, ...);
    for (;;) {
        InterruptWait (0, NULL);
        // do the interrupt work here, at thread priority
        InterruptUnmask (intnum, id);
    }
}
```

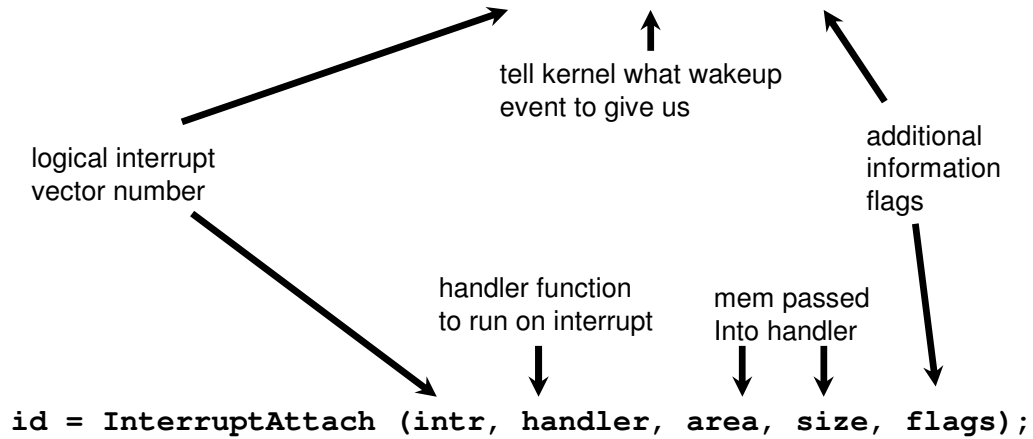


NOTES:

When the kernel gets control, it will mask the interrupt, and use the event to do the appropriate scheduling. In this case, because you are using **SIGEV_INTR**, the *InterruptWait()* will unblock.

Telling kernel what code to run when an interrupt happens:

```
id = InterruptAttachEvent (intr, event, flags);
```



NOTES:

InterruptAttachEvent() and *InterruptAttach()*'s flags parameter can contain:

_NTO_INTR_FLAGS_END: If multiple interrupt handlers, specify we should execute last

_NTO_INTR_FLAGS_PROCESS: for events types that are directed at a process rather than a thread, e.g. pulses

_NTO_INTR_FLAGS_TRK_MSK: request that the kernel adjust the interrupt mask when the attaching process terminates (ALWAYS set this flag)

NOTES:

An interrupt handler operates in the following environment:

- it is sharing the data area of the process that attached it
- the environment is very restricted:
 - cannot call kernel functions except *InterruptMask()*, *InterruptUnmask()* and *TraceEvent()* (see notes)
 - cannot call any function that might call a kernel function
 - the documentation for each function specifies whether or not that function is safe to call from an interrupt handler
 - there is also a section in the Library Reference manual called “Summary of Safety Information” that lists all safe functions
 - the interrupt handler is using the kernel's stack, so keep stack usage small (if you have a lot of data, use variables defined outside of the handler rather than variables defined local to the function, and don't call too many function levels deep)
 - can't do floating point

NOTES:

InterruptLock() and *InterruptUnlock()* may also be called in an interrupt handler as they are not kernel calls. They are implemented as inline assembly.

See the Library Reference for a caveat about *TraceEvent()*.

Should you attach a handler or an event?

- The kernel is the single point of failure for a QNX system; attaching a handler increases the size of the SPOF, an event does not
- debugging is far simpler with an event
 - ISR code can not be stepped/traced with the debugger
- full OS functionality when doing h/w handling in a thread
- events impose far less system overhead at interrupt time than handlers
 - no need for the MMU work to gain access to process address space if using an event
- scheduling a thread for every interrupt could be more overhead, if you could do some work at interrupt time and only need to schedule a thread some of the time
- handlers have lower latency than getting a thread scheduled
 - does your hardware have some sort of buffer or FIFO? If not, then you might not be able to wait until a thread is scheduled

NOTES:

Recommended interrupt event types:

- SIGEV_INTR/*InterruptWait()*
 - simplest to use and fastest
 - must dedicate a thread
 - queue is only 1 entry deep
 - initialize with SIGEV_INTR_INIT()
- Pulse
 - can have multiple threads waiting to receive on the channel
 - are queued
 - most flexible
 - initialize with SIGEV_PULSE_INIT()
- Signal
 - most expensive solution if using a signal handler, but slightly faster than a pulse if waiting with *sigwaitinfo()*
 - can be queued
 - initialize with SIGEV_SIGNAL_INIT() or other signal init macros

NOTES:

Simple interrupt handler:

- In your `interrupt` project is a skeleton file called `intsimple.c`
- Fill it in with the code for handling interrupts, the instructor will tell you which interrupt to attach to
- attach an interrupt handler that will return a **SIGEV_INTR** event
- In the loop, use *InterruptWait()* to wait for the interrupt notification

NOTES:

Topics:

Hardware I/O

Programming PCI bus devices

Handling Interrupts

→ Conclusion

NOTES:

You learned:

- That memory or port mappings have to be set up to access hardware devices
- that the kernel is the first handler for all interrupts
- that processes can register handlers or can register for notification of interrupts
- that interrupt handlers run in a very restricted environment

NOTES: